

CHENLONG ZHANG(张晨龙)

✉ zhangchenlong2023@ia.ac.cn · 🕒 chenlong-clock · 🔍 Google Scholar

🎓 EDUCATION

Institute of Automation, Chinese Academy of Sciences, Beijing, China

2023 – Present

M.S. in NLKE Group, Supervised by Prof. Yubo Chen, GPA: 3.82/4.0

Henan University, Henan, China

2019 – 2023

B.S. in Software Engineering, GPA: 3.84/4.0, Rank: 1/388

💡 RESEARCH INTEREST

Event Understanding and Analysis in LLMs: Exploring how LLMs process and leverage event knowledge to enable applications like timeline summarization, event prediction, and real-world event simulation. **Learning Dynamics of LLMs:** Enhancing the adaptability of language models by addressing catastrophic forgetting (continual learning) and balancing knowledge forgetting and retention (LLM unlearning).

📖 PUBLICATIONS

Chenlong Zhang, Zhuoran Jin, Hongbang Yuan, Jiaheng Wei, Tong Zhou, Kang Liu, Jun Zhao and Yubo Chen. **RULE: Reinforcement Unlearning Achieves Forget-retain Pareto Optimality.** **Under review.**

- **Summary:** We introduce Reinforcement Unlearning (**RULE**) to address unlearning in large language models using boundary optimization with minimal forget data (12%) and synthetic boundary queries (8%).
- **Contributions:** We propose naturalness as a new evaluation dimension for unlearning. **RULE** significantly enhances forgetting quality and response naturalness while robustly generalizing, achieving a forget-retain Pareto optimality against existing methods.
- **Source Codes:** <https://github.com/chenlong-clock/RULE-Unlearn>

Chenlong Zhang, Tong Zhou, Pengfei Cao, Zhuoran Jin, Yubo Chen, Kang Liu and Jun Zhao. **DTELS: Towards Dynamic Granularity of Timeline Summarization.** Accepted in **NAACL 2025 Main**.

- **Summary:** We introduce the task of Dynamic-Granularity Timeline Summarization (**DTELS**) that tailors news timelines to user-specified detail levels and rigorously evaluates them with event-centric metrics.
- **Contributions:** We propose new evaluation metrics based on the nature of events, create the large-scale DTELS-Bench dataset (543 topics, 55k articles, 3 granularity levels), and establish LLM baselines, revealing significant improvements and find that existing LLMs still struggle with capturing fine-grained temporal relationships and consistency across varying levels of granularity.
- **Source Codes:** <https://github.com/chenlong-clock/DTELS-Bench>.

Chenlong Zhang, Pengfei Cao, Yubo Chen, Kang Liu, Zhiqiang Zhang, Mengshu Sun and Jun Zhao. **Continual Few-shot Event Detection via Hierarchical Augmentation Networks.** Accepted in **COLING 2024 Main**.

- **Summary:** We introduce the continual few-shot event detection (**CFED**) paradigm, enabling models to detect events from as few as one labelled instance while preserving performance on previous types.
- **Contributions:** We propose Hierarchical Augmentation Network(**HANet**), combining *prototypical augmentation* for one-exemplar replay with *contrastive augmentation* to curb few-shot over-fitting, yielding up to an 8.4 % micro-F₁ gain over prior methods and even surpassing full retraining and GPT-3.5 baselines.
- **Source Codes:** <https://github.com/chenlong-clock/CFED-HANet>.

👥 PROJECTS

Reference-free Timeline Evaluation via Event Coverage Entropy

2024.11 – 2025.03

Existing timeline annotation requires comprehending thousands of news articles and the consensus on hundreds of events. We present **Reference-free Timeline Evaluation (RTEval)**, an evaluation framework that utilizes “event coverage entropy” to assess the quality of a timeline without annotation.

- We define the quality of a timeline based on the degree to which it can reconstruct the entire topic. The degree of reconstruction can be calculated by “event coverage entropy”.
- We carefully design factual MCQA, applying LLM-as-a-Judge to approximately measure the entropy.

🏆 COMPETITION ORGANIZATION

CCKS 2025 Shared Task: Contextual Event Timeline Generation for Social Media 2025.06 – 2025.08

Designed and co-organized a CCKS 2025 evaluation task targeting event timeline generation from time-stamped, multi-source social media and news documents. The task introduces challenges in temporal event modeling and granularity-controlled summarization. Participants are required to generate structured event timelines satisfying user-defined granularity requirements. We constructed a multi-granularity annotated benchmark dataset, proposed three novel evaluation metrics (informativeness, granularity consistency, factuality), and supported evaluation on real-world event streams including breaking news and public opinion trends.

SemEval 2026 Task: Towards Real-World Event Causal Inference 2026

Proposed and co-authored a new SemEval shared task aiming to benchmark large language models on causal inference over real-world news events. The task, Abductive Event Reasoning (AER), focuses on inferring the primary cause behind complex events from retrieved evidence, and highlights challenges in long-context modeling, semantic confusion, and plausible-but-wrong inference. We constructed a high-quality dataset with human annotation (87% consistency), designed multi-choice evaluation metrics, and conducted baseline studies with GPT-4, Qwen2.5, and GLM-4 to assess model performance.

♡ SELECTED HONORS & AWARDS

University of CAS Merit Student, UCAS	2024
Outstanding Graduates of Henan Province, Henan Ministry of Education	Jun. 2023
Excellent Bachelor Thesis, Henan University.	2023
Bluesky Scholarship (top 0.5% in HENU) ×2.	2020 – 2021
National Scholarship (top 0.2% in China) ×3.	2019 – 2023