

CHENLONG ZHANG(张晨龙)

✉ E-mail · 🌐 Github · 🔍 Google Scholar · 🏠 Homepage

🎓 EDUCATION

Institute of Automation, Chinese Academy of Sciences, Beijing, China 2023 – Present

M.S. in Pattern Recognition and Intelligent Systems (GPA: 3.82/4.0)

Supervised by Prof. [Yubo Chen](#) and Prof. [Jun Zhao](#)

Henan University, Henan, China

2019 – 2023

B.S. in Software Engineering (GPA: 3.84/4.0, Rank: 1/388)

</> INTERNSHIPS

The Hong Kong University of Science and Technology (Guangzhou) Feb. 2025 – May. 2025

Research Intern, advised by [Jiaheng Wei](#)

NLPR, Institute of Automation, Chinese Academy of Sciences

Sep. 2022 – Jul. 2023

Research Intern, advised by [Pengfei Cao](#)

💡 RESEARCH INTERESTS

Event Modeling and Simulation with LLMs: How can large language models model complex real-world events, and can LLM simulation on complex events lead to better decision-making?

Trustworthy LLMs: Can we make language models forget specific knowledge and remain robust against adversarial manipulation in sensitive or evolving domains?

Efficient and Adaptable LLMs: Can LLMs continually incorporate new knowledge without catastrophic forgetting? How to efficiently train LLMs under data- and supervision-limited conditions?

📖 PUBLICATIONS

Chenlong Zhang, Zhuoran Jin, Hongbang Yuan, Jiaheng Wei, Tong Zhou, Kang Liu, Jun Zhao and Yubo Chen. RULE: Reinforcement Unlearning Achieves Forget-retain Pareto Optimality. **Under review.**

- **Summary:** We introduce Reinforcement Unlearning (**RULE**) to address unlearning using boundary optimization with minimal forget data (12%) and synthetic boundary queries (8%).
- **Contributions:** We propose naturalness as a new evaluation dimension for unlearning. RULE significantly enhances forgetting quality and response naturalness while robustly generalizing, achieving a forget-retain Pareto optimality against existing methods.
- **Source codes:** <https://github.com/chenlong-clock/RULE-Unlearn>

Chenlong Zhang, Tong Zhou, Pengfei Cao, Zhuoran Jin, Yubo Chen, Kang Liu and Jun Zhao. **DTELS: Towards Dynamic Granularity of Timeline Summarization.** Accepted in **NAACL 2025 Main**.

- **Summary:** We introduce Dynamic-Granularity Timeline Summarization (**DTELS**) task that tailors timelines to user-specified levels and rigorously evaluates them with event-centric metrics.
- **Contributions:** We propose new evaluation metrics based on the nature of events, create the large-scale DTELS-Bench dataset (543 topics, 55k articles, 3 granularity levels), and establish LLM baselines, revealing significant improvements and find that existing LLMs still struggle with capturing fine-grained temporal relationships and consistency across varying levels of granularity.
- **Source codes:** <https://github.com/chenlong-clock/DTELS-Bench>.

Chenlong Zhang, Pengfei Cao, Yubo Chen, Kang Liu, Zhiqiang Zhang, Mengshu Sun and Jun Zhao. **Continual Few-shot Event Detection via Hierarchical Augmentation Networks.** Accepted in **COLING 2024 Main**.

- **Summary:** We introduce the continual few-shot event detection (**CFED**) paradigm. We propose a rehearsal-based algorithm that enable models to detect events from as few as one labelled instance while preserving performance on previous types.
- **Contributions:** We propose Hierarchical Augmentation Network(**HANet**), combining *prototypical augmentation* for one-exemplar replay with *contrastive augmentation* to curb few-shot overfitting, yielding up to an 8.4 % micro-F₁ gain over prior methods and even surpassing full re-training and GPT-3.5 baselines.
- **Source codes:** <https://github.com/chenlong-clock/CFED-HANet>.

PROJECTS

Reference-free Timeline Evaluation via Event Coverage Entropy 2024.11 – 2025.03

Existing timeline annotation requires comprehending thousands of news articles and the consensus on hundreds of events. We present **Reference-free Timeline Evaluation (RTEval)**, an evaluation framework that utilizes “event coverage entropy” to assess the timeline quality without annotation.

- We define the quality of a timeline based on the degree to which it can reconstruct the entire topic. The degree of reconstruction can be calculated by “event coverage entropy”.
- We carefully design factual MCQA, applying LLM-as-a-Judge to measure the entropy.

HONORS AND AWARDS

Merit Student, University of Chinese Academy of Sciences	May 2024
Outstanding Graduate, Provincial Education Department	Jun. 2023
Excellent Bachelor’s Thesis Award, Henan University	Jun. 2023
China National Scholarship (Top 0.2% nationwide), Ministry of Education	Dec. 2022
Excellent Completion, National Undergraduate Innovation Program (Leader)	Apr. 2022
China National Scholarship (Top 0.2% nationwide), Ministry of Education	Dec. 2021
Bluesky Scholarship (Top 0.5%), Henan University	Jan. 2021
China National Scholarship (Top 0.2% nationwide), Ministry of Education	Dec. 2020

ACADEMIC SERVICES

I actively serve in the NLP community as a shared task organizer and reviewer.

- Task Organizer, **SemEval 2026**: *Abductive Event Reasoning: Towards Real-World Event Causal Inference for Large Language LLMs*
- Task Organizer, **CCKS 2025**: *Event Timeline Generation for Social Media*
- Conference Reviewer: **NLPCC 2025**